

```
In [3]: val df = spark.read
        .format("csv")
        .option("inferSchema", "true")
        .option("header", "true")
        .load("../.../src/201508_station_data.csv")

df.show(5)
```

station_id	name	lat	long	dockcount	landmark	installation
2	San Jose Diridon ...	37.329732	-121.901782	27	San Jose	8/6/2013
3	San Jose Civic Ce...	37.330698	-121.888979	15	San Jose	8/5/2013
4	Santa Clara at Al...	37.333988	-121.894902	11	San Jose	8/6/2013
5	Adobe on Almaden	37.331415	-121.8932	19	San Jose	8/5/2013
6	San Pedro Square	37.336721	-121.894074	15	San Jose	8/7/2013

only showing top 5 rows

Out[3]: df: org.apache.spark.sql.DataFrame = [station_id: int, name: string ... 5 more fields]

```
In [4]: df.printSchema()
```

```
root
|-- station_id: integer (nullable = true)
|-- name: string (nullable = true)
|-- lat: double (nullable = true)
|-- long: double (nullable = true)
|-- dockcount: integer (nullable = true)
|-- landmark: string (nullable = true)
|-- installation: string (nullable = true)
```

Converting to Spark Types

```
In [5]: df.select(lit(1), lit("one"), lit(1.0)).show(2)
```

1	one	1.0
1	one	1.0

only showing top 2 rows

```
In [6]: df.select(lit(1), lit("one"), lit(1.0)).printSchema()
```

```
root
|-- 1: integer (nullable = false)
|-- one: string (nullable = false)
|-- 1.0: double (nullable = false)
```

```
In [7]: df.withColumn("one", lit(1)).show(2)
```

station_id	name	lat	long	dockcount	landmark	installation	one
2	San Jose Diridon ...	37.329732	-121.901782	27	San Jose	8/6/2013	1
3	San Jose Civic Ce...	37.330698	-121.888979	15	San Jose	8/5/2013	1

only showing top 2 rows

Working with Booleans

```
In [11]: df.createOrReplaceTempView("df")
spark.sql("""
SELECT * FROM df
WHERE dockcount = 15
""").show(4)

df.where(df("dockcount") == 15).selectExpr("*").show(4)
```

station_id	name	lat	long	dockcount	landmark	installation
3	San Jose Civic Ce...	37.330698	-121.888979	15	San Jose	8/5/2013
6	San Pedro Square	37.336721	-121.894074	15	San Jose	8/7/2013
7	Paseo de San Antonio	37.333798	-121.886943	15	San Jose	8/7/2013
8	San Salvador at 1st	37.330165	-121.885831	15	San Jose	8/5/2013

only showing top 4 rows

```
In [8]: spark.sql("""
SELECT * FROM df
WHERE dockcount <> 27
""").show(4)

df.where(col("dockcount") != 27).selectExpr("*").show(4)
```

station_id	name	lat	long	dockcount	landmark	installation
3	San Jose Civic Ce...	37.330698	-121.888979	15	San Jose	8/5/2013
4	Santa Clara at Al...	37.333988	-121.894902	11	San Jose	8/6/2013
5	Adobe on Almaden	37.331415	-121.8932	19	San Jose	8/5/2013
6	San Pedro Square	37.336721	-121.894074	15	San Jose	8/7/2013

only showing top 4 rows

```
In [9]: spark.sql("""
SELECT * FROM df
WHERE landmark <> 'San Jose'
""").show(4)

df.where("landmark <> 'San Jose'").selectExpr("*").show(4)
```

station_id	name	lat	long	dockcount	landmark	installation
21	Franklin at Maple	37.481758	-122.226904	15	Redwood City	8/12/2013
22	Redwood City Calt...	37.486078	-122.232089	25	Redwood City	8/15/2013
23	San Mateo County ...	37.487616	-122.229951	15	Redwood City	8/15/2013
24	Redwood City Publ...	37.484219	-122.227424	15	Redwood City	8/12/2013

only showing top 4 rows

```
In [20]: spark.sql("""
SELECT * FROM df
WHERE landmark = 'San Francisco'
""").show(4)

df.where("landmark = 'San Francisco'").selectExpr("*").show(4)
```

station_id	name	lat	long	dockcount	landmark	installation
41	Clay at Battery	37.795001	-122.39997	15	San Francisco	8/19/2013
42	Davis at Jackson	37.79728	-122.398436	15	San Francisco	8/19/2013
45	Commercial at Mon...	37.794231	-122.402923	15	San Francisco	8/19/2013
46	Washington at Kea...	37.795425	-122.404767	15	San Francisco	8/19/2013

only showing top 4 rows

```
In [12]: val stationFilter = col("station_id") > 6
val landmarkFilter = col("landmark").contains("San")

df.where(stationFilter.or(landmarkFilter)).show(15)
```

station_id	name	lat	long	dockcount	landmark	installation
2	San Jose Diridon ...	37.329732	-121.901782	27	San Jose	8/6/2013
3	San Jose Civic Ce...	37.330698	-121.888979	15	San Jose	8/5/2013
4	Santa Clara at Al...	37.333988	-121.894902	11	San Jose	8/6/2013
5	Adobe on Almaden	37.331415	-121.8932	19	San Jose	8/5/2013
6	San Pedro Square	37.336721	-121.894074	15	San Jose	8/7/2013
7	Paseo de San Antonio	37.333798	-121.886943	15	San Jose	8/7/2013
8	San Salvador at 1st	37.330165	-121.885831	15	San Jose	8/5/2013
9	Japantown	37.348742	-121.894715	15	San Jose	8/5/2013
10	San Jose City Hall	37.337391	-121.886995	15	San Jose	8/6/2013
11	MLK Library	37.335885	-121.88566	19	San Jose	8/6/2013
12	SJSU 4th at San C...	37.332808	-121.883891	19	San Jose	8/7/2013
13	St James Park	37.339301	-121.889937	15	San Jose	8/6/2013
14	Arena Green / SAP...	37.332692	-121.900084	19	San Jose	8/5/2013
16	SJSU - San Salvad...	37.333955	-121.877349	15	San Jose	8/7/2013
21	Franklin at Maple	37.481758	-122.226904	15	Redwood City	8/12/2013

only showing top 15 rows

```
Out[12]: stationFilter: org.apache.spark.sql.Column = (station_id > 6)
landmarkFilter: org.apache.spark.sql.Column = contains(landmark, San)
```

```
In [14]: df.where(stationFilter).where(landmarkFilter).show(10)

// same thing with
df.where(stationFilter.and(landmarkFilter)).show(10)
```

station_id	name	lat	long	dockcount	landmark	installation
7	Paseo de San Antonio	37.333798	-121.886943	15	San Jose	8/7/2013
8	San Salvador at 1st	37.330165	-121.885831	15	San Jose	8/5/2013
9	Japantown	37.348742	-121.894715	15	San Jose	8/5/2013
10	San Jose City Hall	37.337391	-121.886995	15	San Jose	8/6/2013
11	MLK Library	37.335885	-121.88566	19	San Jose	8/6/2013
12	SJSU 4th at San C...	37.332808	-121.883891	19	San Jose	8/7/2013
13	St James Park	37.339301	-121.889937	15	San Jose	8/6/2013
14	Arena Green / SAP...	37.332692	-121.900084	19	San Jose	8/5/2013
16	SJSU - San Salvad...	37.333955	-121.877349	15	San Jose	8/7/2013
41	Clay at Battery	37.795001	-122.39997	15	San Francisco	8/19/2013

only showing top 10 rows

Working with Numbers

```
In [21]: val test = (pow(col("station_id"), 3) + 3) * col("dockcount")
df.select(expr("*"), test.alias("test")).show(4)

df.selectExpr("*", "(POWER(station_id, 3) + 3) * dockcount as test").show(4)

spark.sql("""
SELECT *, ((POWER(station_id, 3) + 3) * dockcount) AS test
FROM df """).show(4)
```

station_id	name	lat	long	dockcount	landmark	installation	test
2	San Jose Diridon ...	37.329732	-121.901782	27	San Jose	8/6/2013	297.0
3	San Jose Civic Ce...	37.330698	-121.888979	15	San Jose	8/5/2013	450.0
4	Santa Clara at Al...	37.333988	-121.894902	11	San Jose	8/6/2013	737.0
5	Adobe on Almaden	37.331415	-121.8932	19	San Jose	8/5/2013	2432.0

only showing top 4 rows

```
In [13]: df.withColumn("manyDock", not(df("dockcount").leq(25)))
        .filter("manyDock")
        .select("station_id", "name", "dockcount").show(3)

df.withColumn("notManyDock", expr("NOT dockcount >= 12"))
        .filter("notManyDock")
        .select("station_id", "name", "dockcount").show(3)
```

station_id	name	dockcount
2	San Jose Diridon ...	27
61	2nd at Townsend	27
67	Market at 10th	27

only showing top 3 rows

station_id	name	dockcount
4	Santa Clara at Al...	11
32	Castro Street and...	11
35	University and Em...	11

only showing top 3 rows

```
In [16]: spark.sql("""
SELECT round(lat) AS lat_rounded
FROM df """).show(3)

df.select(round(col("lat"), 2).as("lat_rounded")).show(3)
```

lat_rounded
37.0
37.0
37.0

only showing top 3 rows

```
In [14]: df.withColumn("brouned_lat", round(df("lat"))).show(3)
```

station_id	name	lat	long	dockcount	landmark	installation	brouned_lat
2	San Jose Diridon ...	37.329732	-121.901782	27	San Jose	8/6/2013	37.0
3	San Jose Civic Ce...	37.330698	-121.888979	15	San Jose	8/5/2013	37.0
4	Santa Clara at Al...	37.333988	-121.894902	11	San Jose	8/6/2013	37.0

only showing top 3 rows

```
In [17]: df.withColumn("rounded_lat", round(col("lat"), 2)).show(3)
```

station_id	name	lat	long	dockcount	landmark	installation	rounded_lat
2	San Jose Diridon ...	37.329732	-121.901782	27	San Jose	8/6/2013	37.33
3	San Jose Civic Ce...	37.330698	-121.888979	15	San Jose	8/5/2013	37.33
4	Santa Clara at Al...	37.333988	-121.894902	11	San Jose	8/6/2013	37.33

only showing top 3 rows

General stats

```
In [34]: df.select("station_id", "name", "lat", "dockcount").describe().show()
```

summary	station_id	name	lat	dockcount
count	70	70	70	70
mean	43.0	null	37.59024338428572	17.65714285714286
stddev	24.166091947189145	null	0.20347253639672416	4.010441857493954
min	2	2nd at Folsom	37.329732	11
max	84	Yerba Buena Cente...	37.80477	27

Correlation

```
In [37]: df.stat.corr("station_id", "dockcount")

Out[37]: res33: Double = 0.24015841145323474
```

```
In [38]: df.select(corr("station_id", "dockcount")).show()

spark.sql("SELECT corr(station_id, dockcount) FROM df").show()
```

corr(station_id, dockcount)
0.24015841145323474

Crosstab

According to [the official documentation](http://spark.apache.org/docs/latest/api/python/pyspark.sql.html?highlight=crosstab#pyspark.sql.DataFrame.crosstab) (<http://spark.apache.org/docs/latest/api/python/pyspark.sql.html?highlight=crosstab#pyspark.sql.DataFrame.crosstab>) the crosstab method computes a pair-wise frequency table of the given columns. Also known as a contingency table. The number of distinct values for each column should be less than 1e4. At most 1e6 non-zero pair frequencies will be returned. The first column of each row will be the distinct values of col1 and the column names will be the distinct values of col2.

```
In [15]: //df.stat.crosstab("lat", "long").show()
```

Whereas [freqItems](http://spark.apache.org/docs/latest/api/python/pyspark.sql.html?highlight=freqItems#pyspark.sql.DataFrame.freqItems) (<http://spark.apache.org/docs/latest/api/python/pyspark.sql.html?highlight=freqItems#pyspark.sql.DataFrame.freqItems>) let's you find frequent items for columns, possibly with false positives.

```
In [41]: df.stat.freqItems(Seq("lat", "long")).show()
```

lat_freqItems	long_freqItems
[37.80477, 37.798...]	[-122.229951, -12...]

Adding a unique ID to each row (note that the number doesn't necessarily follow each other)

```
In [45]: df.select(monotonically_increasing_id(), col("station_id")).show(5)
```

monotonically_increasing_id()	station_id
0	2
1	3
2	4
3	5
4	6

only showing top 5 rows

Working with Strings

```
In [46]: df.show(2)
```

station_id	name	lat	long	dockcount	landmark	installation
2	San Jose Diridon ...	37.329732	-121.901782	27	San Jose	8/6/2013
3	San Jose Civic Ce...	37.330698	-121.888979	15	San Jose	8/5/2013

only showing top 2 rows

```
In [51]: df.select(lower(col("Name")), upper(col("Name")), initcap(lower(col("Name")))).show(2, false)

spark.sql("SELECT lower(Name), upper(Name), initcap(lower(Name)) FROM df").show(2)
```

lower(Name)	upper(Name)	initcap(lower(Name))
san jose diridon caltrain station	SAN JOSE DIRIDON CALTRAIN STATION	San Jose Diridon Caltrain Station
san jose civic center	SAN JOSE CIVIC CENTER	San Jose Civic Center

only showing top 2 rows

Note that if lpad or rpad takes a number less than the length of the string, it will always remove values from the right side of the string.

```
In [55]: df.select(
  ltrim(lit("    HELLO    ")).as("ltrim"),
  rtrim(lit("    HELLO    ")).as("rtrim"),
  trim(lit("    HELLO    ")).as("trim"),
  lpad(lit("HELLO"), 3, " ").as("lp"),
  rpad(lit("HELLO"), 10, " ").as("rp")).show(2)

spark.sql("""SELECT
  ltrim('    HELLLLO  '),
  rtrim('    HELLLLO  '),
  trim('    HELLLLO  '),
  lpad('HELLO', 3, ' '),
  rpad('HELLO', 10, ' ')
FROM df""").show(2)
```

ltrim(HELLLLO)	rtrim(HELLLLO)	trim(HELLLLO)	lpad(HELLO, 3,)	rpadd(HELLO, 10,)
HELLLO	HELLLO	HELLLO	HEL	HELLO
HELLLO	HELLLO	HELLLO	HEL	HELLO

only showing top 2 rows

Regular Expressions

There are two key functions in Spark that you'll need in order to perform regular expression tasks: regexp_extract and regexp_replace. These functions extract values and replace values, respectively.

Replace substitute specific names in our name column with TEST:

```
In [36]: val dfBis = df.withColumn("name", upper(col("name")))

import org.apache.spark.sql.functions.regexp_replace

val simpleWords = Seq("san", "santa", "adobe")
val regexString = simpleWords.map(_.toUpperCase).mkString("|")
// the | signifies `OR` in regular expression syntax

dfBis.select(
  regexp_replace(col("name"), regexString, "TEST").alias("name_clean"),
  col("name")).show(5)
```

name_clean	name
TEST JOSE DIRIDON...	SAN JOSE DIRIDON ...
TEST JOSE CIVIC C...	SAN JOSE CIVIC CE...
TESTTA CLARA AT A...	SANTA CLARA AT AL...
TEST ON ALMADEN	ADOBE ON ALMADEN
TEST PEDRO SQUARE	SAN PEDRO SQUARE

only showing top 5 rows

```
Out[36]: dfBis: org.apache.spark.sql.DataFrame = [station_id: int, name: string ... 5 more fields]
import org.apache.spark.sql.functions.regexp_replace
simpleWords: Seq[String] = List(san, santa, adobe)
regexString: String = SAN|SANTA|ADOBE
```

Replace given characters with other characters.

```
In [67]: df.select(translate(col("landmark"), "San", "S4N"), col("landmark")).show(2)
```

translate(landmark, San, S4N)	landmark
S4N Jose	San Jose
S4N Jose	San Jose

only showing top 2 rows

Pulling out the first mentioned word:

```
In [39]: import org.apache.spark.sql.functions.regexp_extract
val regexString = simpleWords.map(_.toUpperCase).mkString("(", "|", ")")
// the | signifies OR in regular expression syntax
dfBis.select(
  regexp_extract(col("name"), regexString, 1).alias("name_clean"),
  col("name")).show(2)
```

name_clean	name
SAN	SAN JOSE DIRIDON ...
SAN	SAN JOSE CIVIC CE...

only showing top 2 rows

```
Out[39]: import org.apache.spark.sql.functions.regexp_extract
regexString: String = (SAN|SANTA|ADOBE)
```

Check for their existence

```
In [40]: dfBis.show(2)
```

station_id	name	lat	long	dockcount	landmark	installation
2	SAN JOSE DIRIDON ...	37.329732	-121.901782	27	San Jose	8/6/2013
3	SAN JOSE CIVIC CE...	37.330698	-121.888979	15	San Jose	8/5/2013

only showing top 2 rows


```
In [43]: val containsSan = col("name").contains("SAN")
val containsJose = col("landmark").contains("Jose")

dfBis.withColumn("test", containsSan.or(containsJose))
      .where("test")
      .select("name", "landmark").show(3, false)

+-----+-----+
|name                |landmark|
+-----+-----+
|SAN JOSE DIRIDON CALTRAIN STATION|San Jose|
|SAN JOSE CIVIC CENTER              |San Jose|
|SANTA CLARA AT ALMADEN              |San Jose|
+-----+-----+
only showing top 3 rows
```

Out[43]: containsSan: org.apache.spark.sql.Column = contains(name, SAN)
containsJose: org.apache.spark.sql.Column = contains(landmark, Jose)

Working with Dates and Timestamps

```
In [23]: val dateDF = spark.range(3)
        .withColumn("current_date", current_date())
        .withColumn("current_timestamp", current_timestamp())

dateDF.createOrReplaceTempView("dateTable")
dateDF.show()
dateDF.printSchema()

+---+-----+-----+
| id|current_date|current_timestamp|
+---+-----+-----+
|  0| 2020-08-29|2020-08-29 15:58:...|
|  1| 2020-08-29|2020-08-29 15:58:...|
|  2| 2020-08-29|2020-08-29 15:58:...|
+---+-----+-----+

root
 |-- id: long (nullable = false)
 |-- current_date: date (nullable = false)
 |-- current_timestamp: timestamp (nullable = false)
```

Out[23]: dateDF: org.apache.spark.sql.DataFrame = [id: bigint, current_date: date ... 1 more field]

```
In [25]: dateDF.select(date_sub(col("current_date"), 5), date_add(col("current_date"), 5)).show(1)
spark.sql("SELECT date_sub(current_date, 5), date_add(current_date, 5) FROM dateTable").show(1)

+-----+-----+
|date_sub(current_date, 5)|date_add(current_date, 5)|
+-----+-----+
|                2020-08-24|                2020-09-03|
+-----+-----+
only showing top 1 row
```

```
In [29]: dateDF.withColumn("week_ago", date_sub(col("current_date"), 7))
        .select(datediff(col("week_ago"), col("current_date"))).show(1)

dateDF.select(
  to_date(lit("2016-01-01")).alias("start"),
  to_date(lit("2017-05-22")).alias("end"))
  .select(months_between(col("start"), col("end"))).show(1)

+-----+
|datediff(week_ago, current_date)|
+-----+
|                -7|
+-----+
only showing top 1 row

+-----+
|months_between(start, end, true)|
+-----+
|                -16.67741935|
+-----+
only showing top 1 row
```

Date doesn't require a format


```
In [30]: spark.range(5).withColumn("date", lit("2017-01-01"))
        .select(to_date(col("date"))).show(1)
```

```
+-----+
|to_date(`date`)|
+-----+
|      2017-01-01|
+-----+
only showing top 1 row
```

The `to_date` function allows you to convert a string to a date, optionally with a specified format. Spark will not throw an error if it cannot parse the date; rather, it will just return null.

```
In [31]: dateDF.select(to_date(lit("2016-20-12")),to_date(lit("2017-12-11"))).show(1)
```

```
+-----+-----+
|to_date('2016-20-12')|to_date('2017-12-11')|
+-----+-----+
|              null|      2017-12-11|
+-----+-----+
only showing top 1 row
```

```
In [110]: val dateFormat = "yyyy-dd-MM"
val cleanDateDF = spark.range(1).select(
  to_date(lit("2017-12-11"), dateFormat).alias("date"),
  to_date(lit("2017-20-12"), dateFormat).alias("date2"))

cleanDateDF.createOrReplaceTempView("dateTable2")
cleanDateDF.show()
```

```
+-----+-----+
|      date|      date2|
+-----+-----+
|2017-11-12|2017-12-20|
+-----+-----+
```

```
Out[110]: dateFormat: String = yyyy-dd-MM
cleanDateDF: org.apache.spark.sql.DataFrame = [date: date, date2: date]
```

```
In [47]: val cleanDF = cleanDateDF
        .withColumn("date", to_timestamp(col("date"), dateFormat))
        .withColumn("date2", to_timestamp(col("date2"), dateFormat))

cleanDF.show()
cleanDF.select(to_timestamp(col("date"), dateFormat)).printSchema()

spark.sql("""SELECT to_timestamp(date, 'yyyy-dd-MM'), to_timestamp(date2, 'yyyy-dd-MM')
FROM dateTable2""").show()
```

```
+-----+-----+
|              date|              date2|
+-----+-----+
|2017-11-12 00:00:00|2017-12-20 00:00:00|
+-----+-----+

root
|-- to_timestamp(`date`, 'yyyy-dd-MM'): timestamp (nullable = true)

+-----+-----+
|to_timestamp(datetable2.`date`, 'yyyy-dd-MM')|to_timestamp(datetable2.`date2`, 'yyyy-dd-MM')|
+-----+-----+
|              2017-11-12 00:00:00|              2017-12-20 00:00:00|
+-----+-----+
```

```
Out[47]: cleanDF: org.apache.spark.sql.DataFrame = [date: timestamp, date2: timestamp]
```

```
In [70]: val dateDF = spark.range(1).withColumn("current_date", current_date())
dateDF.show()

dateDF.filter(col("current_date") <= lit("2017-12-12")).show()
dateDF.filter(col("current_date") >= lit("2017-12-12")).show()
```

```
+---+-----+
| id|current_date|
+---+-----+
|  0|  2020-08-29|
+---+-----+

+---+-----+
| id|current_date|
+---+-----+
+---+-----+

+---+-----+
| id|current_date|
+---+-----+
|  0|  2020-08-29|
+---+-----+
```

```
Out[70]: dateDF: org.apache.spark.sql.DataFrame = [id: bigint, current_date: date]
```

```
In [129]: val dateDF = spark.range(1)
          .withColumn("unix_ts", unix_timestamp(current_timestamp, "MM:dd:yyyy hh:mm:ss"))
          .withColumn("time_stamp", col("unix_ts").cast("timestamp")) // or to_timestamp
          .withColumn("date", to_date(col("time_stamp")))

dateDF.show()
```

```
+---+-----+-----+-----+-----+
| id|  unix_ts|      time_stamp|      date|
+---+-----+-----+-----+-----+
|  0|1598722033|2020-08-29 17:27:13|2020-08-29|
+---+-----+-----+-----+-----+
```

```
Out[129]: dateDF: org.apache.spark.sql.DataFrame = [id: bigint, unix_ts: bigint ... 2 more fields]
```

Working with Nulls in Data

Coalesce The coalesce function allows you to select the first non-null value from a set of columns. If there are no null values, so it simply returns the first column.

```
In [5]: val df = spark.read.format("csv")
        .option("header", "true")
        .option("inferSchema", "true")
        .load("../src/2010-12-01.csv")

df.show(3)
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+
|InvoiceNo|StockCode|      Description|Quantity|      InvoiceDate|UnitPrice|CustomerID|      Country|
+-----+-----+-----+-----+-----+-----+-----+-----+
|  536365|  85123A|WHITE HANGING HEA...|      6|2010-12-01 08:26:00|      2.55|  17850.0|United Kingdom|
|  536365|   71053|WHITE METAL LANTERN|      6|2010-12-01 08:26:00|      3.39|  17850.0|United Kingdom|
|  536365|  84406B|CREAM CUPID HEART...|      8|2010-12-01 08:26:00|      2.75|  17850.0|United Kingdom|
+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 3 rows
```

```
Out[5]: df: org.apache.spark.sql.DataFrame = [InvoiceNo: string, StockCode: string ... 6 more fields]
```

```
In [ ]: df("columnName")           // On a specific DataFrame.
col("columnName")                   // A generic column not yet associated with a DataFrame.
col("columnName.field")             // Extracting a struct field
col("`a.column.with.dots`")         // Escape `.` in column names.
"$columnName"                       // Scala short hand for a named column.
expr("a + 1")                       // A column that is constructed from a parsed SQL Expression.
lit("abc")                          // A column that produces a literal (constant) value.
```

```
In [6]: df.select(coalesce(col("StockCode"), df("CustomerID"))).show(3)
```

```
+-----+
|coalesce(StockCode, CustomerID)|
+-----+
|                               |85123A|
|                               |71053 |
|                               |84406B|
+-----+
only showing top 3 rows
```

```
In [31]: //not working examples
//df.select([count(when(isnan(c), c)).alias(c) for c in df.columns]).show()
//df.select(for (c <- df.columns) yield count(when(isnan(c), c)).alias(c)).show()
```

Nb of NaNs

```
In [7]: println(df.filter(df("Description").isNull).count)
println(df.filter(df("Description") == "").count)
```

```
10
0
```

```
In [8]: df.filter(df("Description").isNull || df("Description") == "").count
```

```
Out[8]: res5: Long = 10
```

```
In [9]: for (c <- df.columns) println(c, df.filter(df("Description").isNull || df("Description") == "").count)
```

```
(InvoiceNo,10)
(StockCode,10)
(Description,10)
(Quantity,10)
(InvoiceDate,10)
(UnitPrice,10)
(CustomerID,10)
(Country,10)
```

How to find duplicated lines

```
In [14]: df.groupBy("InvoiceNo", "StockCode", "Description", "Quantity", "InvoiceDate", "UnitPrice", "CustomerID", "Country")
        .count().show(3)
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+
|InvoiceNo|StockCode|Description|Quantity|InvoiceDate|UnitPrice|CustomerID|Country|
|count|
+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+
|536407|22632|HAND WARMER RED P...|6|2010-12-01 11:34:00|1.85|17850.0|United Kingdom|
1|
|536408|84029E|RED WOOLLY HOTTIE...|4|2010-12-01 11:41:00|3.75|14307.0|United Kingdom|
1|
|536409|22531|MAGIC DRAWING SLA...|1|2010-12-01 11:45:00|0.42|17908.0|United Kingdom|
1|
+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+
only showing top 3 rows
```

```
In [ ]: import pyspark.sql.functions as f
df.groupBy(df.columns)\
  .count()\
  .where(f.col('count') > 1)\
  .select(f.sum('count'))\
  .show()
```

```
In [11]: println(df.distinct().count(), df.count())
```

```
(3064,3108)
```

```
In [13]: var dfB = df.dropDuplicates()
dfB.count()
```

```
Out[13]: dfB: org.apache.spark.sql.Dataset[org.apache.spark.sql.Row] = [InvoiceNo: string, StockCode: string ... 6 more fields]
res10: Long = 3064
```

Remove rows that contain nulls : the default is to drop any row in which any value is null

```
In [17]: dfB.na.drop("any").count() // or drop()
```

```
Out[17]: res14: Long = 1924
```

```
In [18]: dfB.na.drop("all").count()
```

```
Out[18]: res15: Long = 3064
```

Fill all null values in columns of type String

```
In [19]: df.na.fill("All Null values become this string")
```

```
Out[19]: res16: org.apache.spark.sql.DataFrame = [InvoiceNo: string, StockCode: string ... 6 more fields]
```

For integers

```
In [20]: df.na.fill(5, Seq("StockCode", "InvoiceNo"))
```

```
Out[20]: res17: org.apache.spark.sql.DataFrame = [InvoiceNo: string, StockCode: string ... 6 more fields]
```

```
In [21]: val fillColValues = Map("StockCode" -> 5, "Description" -> "No Value")
df.na.fill(fillColValues)
```

```
Out[21]: fillColValues: scala.collection.immutable.Map[String,Any] = Map(StockCode -> 5, Description -> No Value)
res18: org.apache.spark.sql.DataFrame = [InvoiceNo: string, StockCode: string ... 6 more fields]
```

A more flexible option that you can use with more than just null values. (only requirement is that this value be the same type as the original value)

```
In [22]: df.na.replace("Description", Map("" -> "UNKNOWN"))
```

```
Out[22]: res19: org.apache.spark.sql.DataFrame = [InvoiceNo: string, StockCode: string ... 6 more fields]
```